



HyperLogLog in Practice

**Algorithmic Engineering of a State of The Art
Cardinality Estimation Algorithm**

Stefan Heule
Marc Nunkesser
Alexander Hall



Contents

- i. Problem Definition & Applications
- ii. Cardinality Estimation Algorithms
- iii. HyperLogLog and Improvements
- iv. Implementations and Examples



i. Problem Definition & Applications



Problem Description

- [Count-distinct problem](#) or Cardinality Estimation Problem
- Formal Definition:
 - Given a stream of N elements with repetitions
 - Find the number of distinct elements D
 - Using M storage units, where $M \ll D$
- Memory requirement tends to be proportional to the cardinality
- Accuracy depends on the algorithm and the underlying data structure
- Set Membership can be inferred by changes in cardinality



Exact Solution

- Store unique elements in Array Data Structure $\rightarrow A_j = X, j \in [0, n]$
- Advantages:
 - Accuracy
 - Cardinality is the Array length $\rightarrow O(1)$
- Disadvantages:
 - Slow Lookup $\rightarrow O(n)$, insertion time depends on cardinality
 - Required memory scales with cardinality $\rightarrow \sum A_j$
- Possible Improvements:
 - Use lossless compression in order to reduce data size
 - Use Set / Tree data structure (B-tree) with lookup time $O(1)$ / $O(\log n)$



Approximate Solution

- Define:
 - Cardinality $\rightarrow n$
 - Hashing function with uniform distribution $h(x) \rightarrow D, D \in [0, S]$
- Add each hash value to a Set Data Structure
 - Lookup time $\rightarrow O(1)$
 - Cardinality calculation $\rightarrow O(1)$
- Collision probability depends on hash value range
- Accuracy is retained if: $S \gg n$
- Memory requirements are still proportional to cardinality



Problem Domain

- Time-series data streams lead to large, transient datasets
- Not feasible or otherwise necessary to maintain the exact original dataset:
 - Cardinality is unknown
 - Memory constraints
 - Sliding window calculations
- Small errors can be tolerated:
 - Eventual consistency
 - Multiple data sources
- Latency versus Accuracy



Applications: Web Analytics & Security

Problem Definition

- Count unique sessions or pageviews per device/browser for given time periods.
- Detect possible spamming or malicious requests.

Solution Modelling

- Combine session ID and URL → Request ID.
- Count unique Request ID.
- Low cardinality of Request ID and high request rate → throttling?
- Calculation happens over a sliding time window.



Applications: Network Security

Problem Definition

- Detect and possibly block malicious network traffic^[6]
- Very high data rate: 40+ Gb/s
- Packet inter-arrival time must remain low ($\sim 8\text{ns}$)
- Attack detection is not the main role of the router, memory should remain low

Solution Modelling

- Flow is sequence of packets identified by the tuple
(Source-IP, Source-Port, Destination-IP, Destination-Port, Protocol)
- Calculate the total packets using different destination ports over a sliding time window



ii. Cardinality Estimation Algorithms



Concepts & Theory - Part I

- Multiset
 - Modification of Set that allows multiples instances of its elements.
 - Number of instances per element is named *multiplicity*.
 - Cardinality is the number of unique elements.
- Uniform distribution:
 - Probability of random variable within the interval **[a,b]** is independent of the variable.
- K-th Order Statistic
 - Equal to kth-smallest value of a statistical sample
 - First order statistic is the sample minimum
- Probabilistic Algorithms / Data Structures
 - Employ randomness as part of their logic
 - Examples: Bloom Filter, Linear Counting, Skip List, HyperLogLog



Concepts & Theory - Part II

- Estimator
 - Rule for calculating an estimate of a given quantity based on observed data
- Bias
 - Expected relative error of an estimator
 - **bias = 0.01** means an overestimation of **1%** compared to the exact value
- Standard Error
 - Standard deviation of the ratio of estimated quantity compared to the exact value
- Harmonic Mean
 - Type of average function, better suited for average of rates
 - Expressed as the reciprocal of the arithmetic mean of the reciprocals

$$H = (\sum x_i^{-1} / n)^{-1}$$



Cardinality Observables - Part I

- Consider the uniformly distributed hashed values of a Multiset (S-values)
- **Bit-pattern Observables**
 - Series of bits (bit strings) of predefined length (e.g. 32-bit)
 - Pattern of bits occurring at the beginning of the S-values
 - Observe the minimum index of a 1-bit
 - By convention we use **LSB order**

Cardinality Observables - Part II

- Assume hashing function of 4 bits, LSB order in bit strings:
 - **50%** of hashed values will display the pattern **1BBB**
 - **25%** of hashed values will display the pattern **01BB**
- Hashing **4** unique elements, it is likely that at least one S-value will start with **01**
- Inverting the expectation:
 - If the first 1-bit index is **2**, it is likely we encountered **~4 unique values**.

ABC	→	14	→	0	1	1	1
DEF	→	11	→	1	1	0	1
GKF	→	15	→	1	1	1	1
PKM	→	7	→	1	1	1	0



Cardinality Observables - Part II

- **Order-Statistic Observables**
 - Hash function provides uniformly distributed real values in **[0, 1]**
 - The minimum value in an ideal multiset does not depend on:
 - Replication structure of the data
 - Ordering
 - Indication on the number of distinct values of the multiset
 - The minimum of independent uniform values on $[0, 1]$ has more chances of being small if cardinality is large
 - MinCount^[9]



Available Probabilistic Algorithms

- [Flajolet - Martin](#) (1985)
- [Linear Counting](#) (Whang et al - 1990)
 - Bit-map of size m - initialized to zero values
 - Hash function generates bit map address which is set to one
 - Cardinality is computed by: $n = -m * \ln(\text{ZeroCount} / m)$
- [LogLog](#) (Durand & Flajolet - 2003)
 - Improvements to the Flajolet - Martin algorithm
- [HyperLogLog](#) (Flajolet et al - 2007)

Probabilistic Counting - An Example

- [Approximate Counting](#) (Morris - 1977)
 - Count large number of events using small amount of memory
 - Counts powers of two - only stores the exponent
 - Incremented by Pseudo-random event: “coin flip” times the current counter value

0 → $0 = 0$

1 → $2^1 = 2$

1 → $2^2 = 4$

0 0 **1 0** **0 1** **1 1** → $2^3 = 8$



iii. HyperLogLog and Improvements



Introduction to HyperLogLog - Part I

- Computational Algorithm presented in [Flajolet et al](#) (2007)
- Contains the analytical mathematical proofs:
 - The estimator is asymptotically unbiased.
 - The standard error constant values used for bias correction.
- Applies hash function on the original data stream
- Uses bit-pattern observables
- Divides the hash values input stream into m substreams
 - Specific bit interval is used for partitioning
- Emulates m experiments in parallel with a single hash function
 - As opposed to using multiple hash values per original input element
- Estimate of the cardinality is a suitable average of the substream observables
 - Quality should improve due to averaging effects as substreams increase
- Solution is named **Stochastic Averaging**



Introduction to HyperLogLog - Part II

- If each of a m random variables has standard deviation σ
 - Their arithmetic mean has deviation σ / \sqrt{m}
- Accuracy Characteristics ($m = 2048$, hashing to 32 bit values):
 - Cardinalities close to and exceeding one billion (10^9)
 - Typical accuracy of 2% error $\rightarrow 1.04 / \sqrt{m}$
 - 1.5 kilobyte of storage
- Substream counters are also referred as registers
- Use of harmonic mean since it is less sensitive to outliers.
- Suggests a correction formula depending on estimate range (E):
 - Small Range ($\leq 2.5m$): Use Linear Counting
 - Large Range ($> 2^{32} / 30$)

HyperLogLog - Computational Algorithm (I)

- Let $h : D \rightarrow [0, 1] \equiv \{0, 1\}^\infty$ hash data from domain D to the binary domain.
- Let $\rho(s)$, for $s \in \{0, 1\}^\infty$, be the position of the leftmost 1-bit ($\rho(0001 \dots) = 4$).
- Algorithm HYPERLOGLOG (input M : multiset of items from domain D).
- Assume $m = 2^b$ with $b \in \mathbb{Z} > 0$
- Initialize a collection of m registers, $M[1], \dots, M[m]$, using a value of $-\infty$
- for v in M :
 - $x \leftarrow h(v)$
 - $j \leftarrow 1 + x_1 x_2 \dots x_b$
 - $w \leftarrow x_{b+1} x_{b+2} \dots$
 - $M[j] = \max(M[j], \rho(w))$
- Compute indicator function $Z = (\sum 2^{-M[j]})^{-1}$
- Compute estimator $E = \alpha_m m^2 Z$ with α_m as a bias-correction constant.

HyperLogLog - Computational Algorithm (II)

- Assume hashing function h with 8-bit output range and 4 registers ($m = 2^2$, $b = 2$)
- Let $v = \text{'abcdef'}$
- $x = h(v) = 142 = 01110001$
- $j = \text{'01110001'}_{1..2} = \text{'01'} = 2$
- $w = \text{'01110001'}_{3..8} = \text{'110001'}$
- $\rho(w) = \rho(\text{'110001'}) = 1$
- $M[2] = 1$
- $x = h(\text{'ghij'}) = 178 = 01001101$
- $j = 2$, $\rho = 3$, $M[2] = \max(M[2], 3)$, $M[2] = 3$



Example Implementation - Part I

```
def compute_hash(value):
    return zlib.crc32(value.encode()) % (1 << 32)

def to_binary_representation(integer):
    return '{0:b}'.format(integer)

def to_lsb_binary(integer):
    return to_binary_representation(integer)[::-1]

def convert_to_base_10(value):
    return int(value, 2)
```

```
def lsb_bit_range(number, start, end):
    return to_lsb_binary(number)[start:end]

def indicator_function(registers):
    return 1. / sum(2 ** -register for
        register in registers)

def find_first_one_bit_index(lsb_binary):
    return lsb_binary.index('1') + 1
```



Example Implementation - Part II

```
def add(registers, b, value):
    x = compute_hash(value)
    j = convert_to_base_10(lsb_bit_range(x, 0, b)[::-1])
    w = lsb_bit_range(x, b, -1)
    rho = find_first_one_bit_index(w)
    changed = rho > registers[j]
    registers[j] = max(rho, registers[j])
    return changed

def count(registers):
    a_m = 0.72134
    Z = indicator_function(registers)
    return a_m * (m ** 2) * Z

def merge(h1l_1, h1l_2):
    return map(max, h1l_1, h1l_2)
```




Example Implementation - Part III

```
def initialize_hll(b):  
    m = 2 ** b  
    return [-math.inf] * m
```

```
B = 11
```

```
hll_registers = initialize_hll(B)
```

```
for n in range(1, 10**7 + 1):  
    if add(registers=hll_registers b=B, value=str(n)):  
        cardinality = count(hll_registers)  
        if n % 100 == 0:  
            print(  
                'n={}, hyperloglog count={}, error={}'.format(  
                    n, cardinality, abs(n - cardinality) / n)  
            )
```



HyperLogLog - Operations and Complexity

- Add
 - Time Complexity $O(1)$
- Merge
 - $HLL_{union}[j] = \text{Max}(HLL_1[j], HLL_2[j])$
 - Dependent on register count: Complexity $O(m)$
- Count
 - Time Complexity $O(m)$
- Space
 - $O(m)$
 - Total size depends on the hashing function output range S
 - Register Size $\log_2(S\text{-bits}) - \log_2 m$



HyperLogLog++ - Part I

- HyperLogLog weaknesses:
 - As cardinality approaches 2^L , where L is the number of hash output bits, number of collisions increases.
 - Zero cardinality for $n \ll m \log m$
- Improvements to the original HyperLogLog algorithm
 - Use a 64-bit hashing function, instead of 32-bit
 - Initialize registers to zero to avoid zero cardinality for $n \ll m \log m$
 - Large range correction no longer necessary due to the 64-bit range shift
 - Empirical bias correction
 - Sparse/dense representation
- Improved Accuracy for cardinalities larger than 2^{32}
- Small increase in memory requirements



HyperLogLog++ - Part II

- **Small Cardinality Estimation**
 - Linear Counting below 2.5m
 - Empirical bias correction until 5m
 - Calculation of the mean difference of raw estimate minus the cardinality
 - Use [k-nearest neighbor interpolation](#)
(200 cardinalities as interpolation points)
- **Sparse Representation**
 - Stores pairs (index, $\rho(w)$) with size threshold 6m bits
 - Represented as sorted list of integers, by concatenating bit patterns:
 <INDEX><RHO>
 - Variable-Length Encoding
 - Difference Encoding
 - Store consecutive differences for the second element and so forth



Sliding HyperLogLog

- Problem: allow calculations over a sliding time window for unknown intervals
- Store a List of Possible Future Maxima for each register
- LPFM entries [Timestamp, R_i]
 - R_i is the $\rho(w)$ value of the HyperLogLog algorithm.
 - Timestamp is the Unix epoch time (seconds)
- Requires approximately 5 additional bytes per LPFM element.
- Query: estimate cardinality given an interval $[t - w, t]$
 - Retrieve all LPFM entries for which the timestamp is within the interval.
 - Maintain the maximum R_i value for each LPFM.
 - Apply stochastic averaging on the computed maximum register values.
- Remove LPFM entries based on the timestamp as the window advances



iv. Implementations and Examples



Elasticsearch

- [Elasticsearch Cardinality Aggregations](#)
- Cardinality Aggregation uses HyperLogLog++ (based on the Google paper).
- Example: Count unique IPs

POST /logs/_search?size=0

```
{
  "aggs" : {
    "type_count" : {
      "cardinality" : {
        "field" : "source_ipv6"
      }
    }
  }
}
```



Elasticsearch

- Configurable precision with the `precision_threshold` option
- Increasing the threshold requires more memory, in order to improve accuracy.
- The threshold defines a unique count below which counts are expected to be close to accurate.
 - Maximum supported value is 40000, with a default value of 3000.
 - Memory: $\text{Threshold} * 8 \text{ bytes}$
- Pre-computed hashes:
 - Value hashing can be performed using an ES plugin.
 - Beneficial for large string and high-cardinality fields.
 - Field hash value is stored in the document.
 - Aggregation is applied on the hash field.



Redis - Part I

- Redis is an in-memory data structure store
- HyperLogLog structures are available since version **2.8.9**^[10]
- Largely based on HyperLogLog++ with further improvements
- **Sparse representation**, optimized for storing large number of registers set to zero:
 - Lossless compression with run-length encoding
- **Dense representation:**
 - Redis string of 12288 bytes in order to store **16384 6-bit counters**
- 6-bit registers
 - With 64-bit hash values ($m = 2^{14}$) 50 bits remain that require 2^6 storage bits
- Maximum memory is **12 kB**, standard error **0.81%** (since 16384 registers are used)
- 64-bit hash function ([MurmurHash2](#))



Redis - Part II

- Introduced curve fitting resulting in a four-order polynomial for error correction in range 40960-72000
 - Linear counting:
$$\text{Cardinality} = m * \ln(m / \text{total-registers-with-initialization-value})$$
 - Finally using [τ raw estimator](#) (Ertl - 2017)
- [PFADD](#) key element [element ...]
 - Adds the given elements to the HLL structure.
 - Returns 1 if any of the internal registers was modified.
- [PFCOUNT](#) key [key]
 - Returns the approximated cardinality of the specified HLL or the combined cardinality by merging multiple HLL structures.
- [PFMERGE](#) destkey source-key [source-key ...]
 - Merge multiple HLL structures to a new one.



Druid

- Column-oriented, distributed data store
 - High performance analytics data store for event-driven data
- [Cardinality Aggregator](#)
- [Fast, Cheap, and 98% Right: Cardinality Estimation for Big Data](#)
 - Murmur 128 hashing function
 - Stores intermediate HLL format in a column
- [How We Scaled HyperLogLog: Three Real-World Optimizations](#)
 - Register compaction:
 - Offset + positive differences in the registers
 - Faster cardinality calculations was to use lookups for register values
 - Dense/sparse storage



Libraries

PostgreSQL extension adding HLL as native datatype:

<https://github.com/citusdata/postgresql-hll>

C#: <https://github.com/Microsoft/CardinalityEstimation>

Golang: <https://github.com/axiomhq/hyperloglog>

Python: <https://github.com/ekzhu/datasketch>

Erlang: <https://github.com/GameAnalytics/hyper>



References

- [1] [HyperLogLog in Practice: Algorithmic Engineering of a State of The Art Cardinality Estimation Algorithm](#)
- [2] [Probabilistic Counting Algorithms for Data Base Applications](#)
- [3] [Streaming Algorithm](#)
- [4] [Elasticsearch: The Definitive Guide \[2.x\] » Aggregations » Approximate Aggregations » Finding Distinct Counts](#)
- [5] [Approximate counting algorithm](#)
- [6] [How can sliding HyperLogLog and EWMA detect port scan attacks in IP traffic?](#)
- [7] [Sliding HyperLogLog: Estimating cardinality in a data stream](#)
- [8] [HyperLogLog: the analysis of a near-optimal cardinality estimation algorithm](#)
- [9] [Order statistics and estimating cardinalities of massive data sets](#)
- [10] [Redis new data structure: HyperLogLog](#)